# GALX

Sustainable, Abundant AI Access
for Anyone Who Wants to Build

**Two** Models

**<100ms** TTFT

**99.99%** Continuity

**Zero** Ops

GALX / A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD

# THE STRUCTURAL INEFFICIENCY: THE STATUS QUO TAX

ONE INTERFACE. MANY WORLDS. REAL PROGRESS.

## UNSHIELDED STATUS QUO

14 APIs. COUNTLESS PROBLEMS.

INTEGRATION
MAINTENANCE
TROUBLESHOOTING
CONSTANT CHANGES
LOST PRODUCTIVITY

FROM CHAOS TO CLARITY

## THE GALX ABSTRACTION SHIELD

ONE INTERFACE. A MORE CAPABLE WORLD.

**68%**
**OPS OVERHEAD**
Average engineering time spent on integration, maintenance, and firefighting.

**14**
**FRAGMENTED SDKs**
Different APIs. Different rules. A moving target.

**<100ms**
**AUTONOMOUS ROUTING**
Intelligent model selection. Optimal performance, automatically.

GALX | UNIFIED INTELLIGENCE REAL PROGRESS.

# THE SUFFICIENCY PRINCIPLE:
## SUFFICIENT INTELLIGENCE ECONOMICALLY DELIVERED

SAME IDEAS.
MORE EXECUTION.
A MORE CAPABLE WORLD.

## TASK COMPLETION vs. MODEL COMPUTE SCALE

Diminishing returns are real. Sufficient intelligence wins.

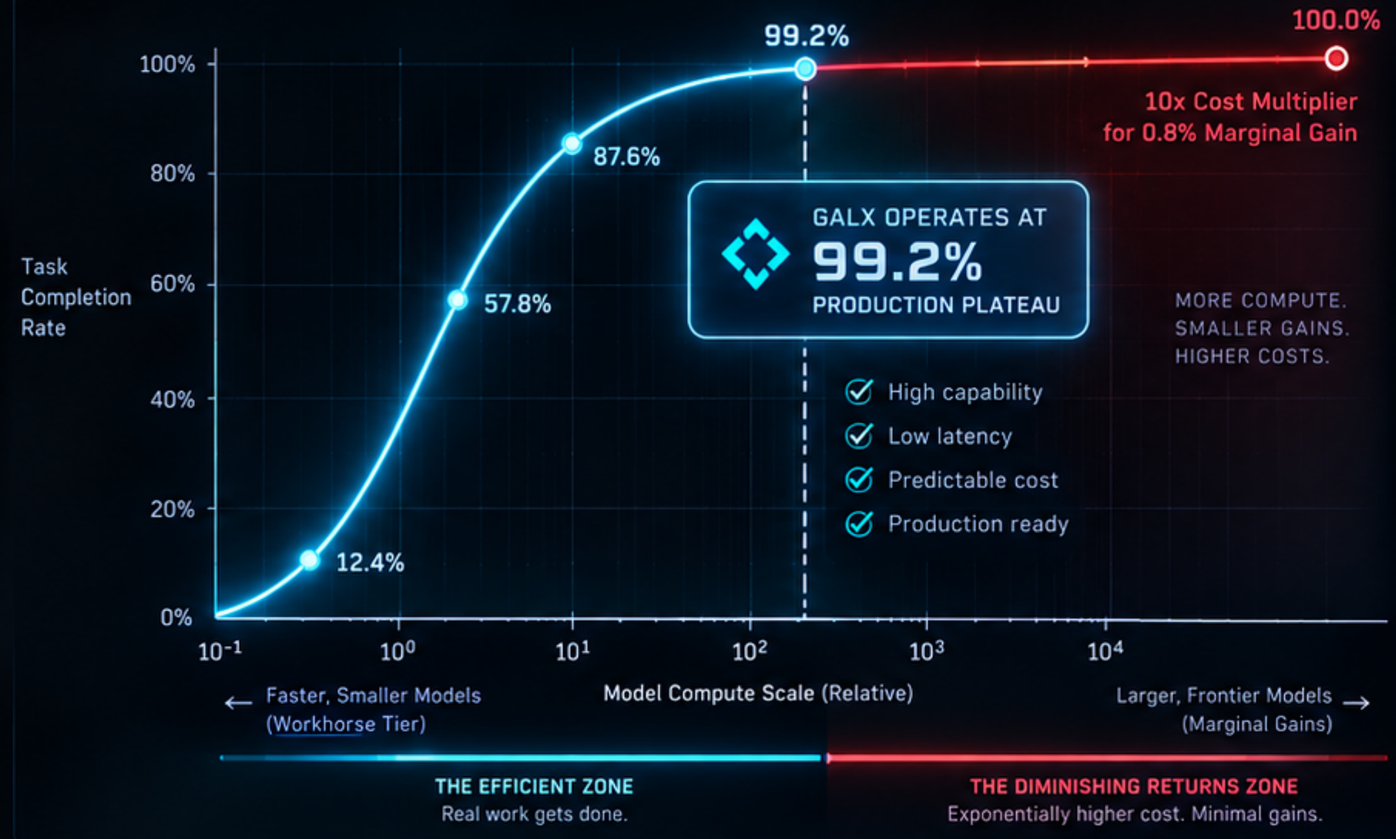


## THE VENTURE LANDSCAPE

From Fragile Sprawl to Sovereign Sufficiency.

**99.2%**
**TASK PLATEAU**
Real-world tasks completed with workhorse models.

**10x**
**COST MULTIPLIER AVOIDED**
Frontier-scale costs for only 0.8% marginal gain.

**2**
**CURATED MODELS**
Sufficient. Capable. Production ready.
No catalog sprawl.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

# RADICAL SIMPLICITY:
## TWO MANAGED MODELS. STANDARD API.

LESS COMPLEXITY.
MORE POSSIBILITY.
SAME INTERFACE.

### THE RIGHT MODEL FOR THE JOB.
PRODUCTION-READY. ECONOMICALLY DELIVERED.

## GALX FAST
**SPEED FOR REAL WORK**

High-throughput, low-latency inference for everyday tasks.
Get to answers, faster.

- **Sub-100ms latency** — Blazing fast responses
- **High throughput** — Built for production scale
- **Micro-cents pricing** — Maximum efficiency
- **Reliable & deterministic** — Production ready

GALX FAST

FAST. EFFICIENT. RELIABLE.

## GALX HEAVY
**DEPTH FOR COMPLEXITY**

Advanced reasoning, long-context understanding, and multi-step agentic workflows.

- **Deep reasoning** — Tackle complex problems
- **High-context synthesis** — Larger context, richer understanding
- **Complex agents** — Multi-step task execution
- **Production-grade reliability** — Same simple interface

GALX HEAVY

DEEPER. SMARTER. FURTHER.

GALX | Code | Response | Logs | Live | cURL

```
curl https://api.galx.ai/v1/chat/completions \
  -H "Content-Type: application/json" \
  -H "Authorization: Bearer $GALX_API_KEY" \
  -d '{
    "model": "galx-fast",
    "messages": [{"role": "user", "content": "Build something amazing."}]
}'
```

Streaming response... 342 tokens/s

```
{
  "id": "galx-9f3c2e1b",
  "object": "chat.completion.chunk",
  "model": "galx-fast",
  "choices": [{
    "delta": {
      "content": "Building",
    }
  }]
}
{  "delta": { "content": " something" } }
{  "delta": { "content": " amazing" } }
{  "delta": { "content": "." } }
```

Status **200 OK**

TTFT **84ms**

Continuity **100%**

**Single Unified Endpoint**

SAME INTERFACE.
TWO MODELS.
INFINITE POTENTIAL.

**1**
**STANDARD ENDPOINT**
One API. Two managed models.
No complexity.

**<100ms**
**TTFT**
Get to answers, faster.
Built for real work.

**ZERO**
**MIGRATION CONFIG**
No SDK sprawl. No rewrites.
Start building today.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

# THE BUILDER WEDGE: HIGH-VELOCITY CREATORS

LESS SETUP. MORE BUILDING. A MORE CAPABLE WORLD.

IDEAS SHIP FASTER HERE.

SAME AMBITION.
A SHORTER PATH.

- IDEA TO PROTOTYPE
- PROTOTYPE TO PRODUCT
- BUILD TO IMPACT

INFRASTRUCTURE SHOULD ACCELERATE HUMAN CREATIVITY.

Vercel | supabase | Cloudflare | Next.js

WORKS WITH YOUR STACK.

## TIME-TO-PRODUCTION

SAME IDEA. TWO VERY DIFFERENT PATHS.

**TRADITIONAL CLOUD / MULTI-VENDOR**

Multiple providers.
Multiple roadblocks.
Weeks of friction.

**4 to 6 Weeks**

SLOWER BUILDS.
HIGHER COSTS.
MORE COMPLEXITY.

**GALX PLATFORM**

One interface.
Instant access.
Built for builders.

**3 Minutes**
**to First Production Token**

SAME IDEAS.
REAL MOMENTUM.

"INFRASTRUCTURE DISAPPEARS. CREATIVITY ACCELERATES."

### 3-MINUTE ONBOARDING

Go from sign up to production token in 3 minutes.

### 4–6 WEEKS TIME SAVED

Skip procurement, integration, and infrastructure complexity.

### 100% FOCUS ON PRODUCT

Spend your time building, not juggling infrastructure.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

# WHY THIS CAN BECOME ENORMOUS: THE COMPOUNDING EXECUTION FLYWHEEL

MORE USAGE. HIGHER EFFICIENCY. STRONGER MARGINS. A LARGER GALX.

SAME INTELLIGENCE. A BIGGER TOMORROW.

DEMAND SCALES INTELLIGENCE. EFFICIENCY COMPOUNDS VALUE.

FROM IDEAS TO IMPACT AT GLOBAL SCALE.

A MORE CAPABLE WORLD IS A COMPOUNDING ONE.

## QUERY VOLUME & GROSS MARGIN EXPANSION

SCALING USAGE DRIVES EFFICIENCY. EFFICIENCY DRIVES MARGINS. MARGINS ENABLE FURTHER SCALE.

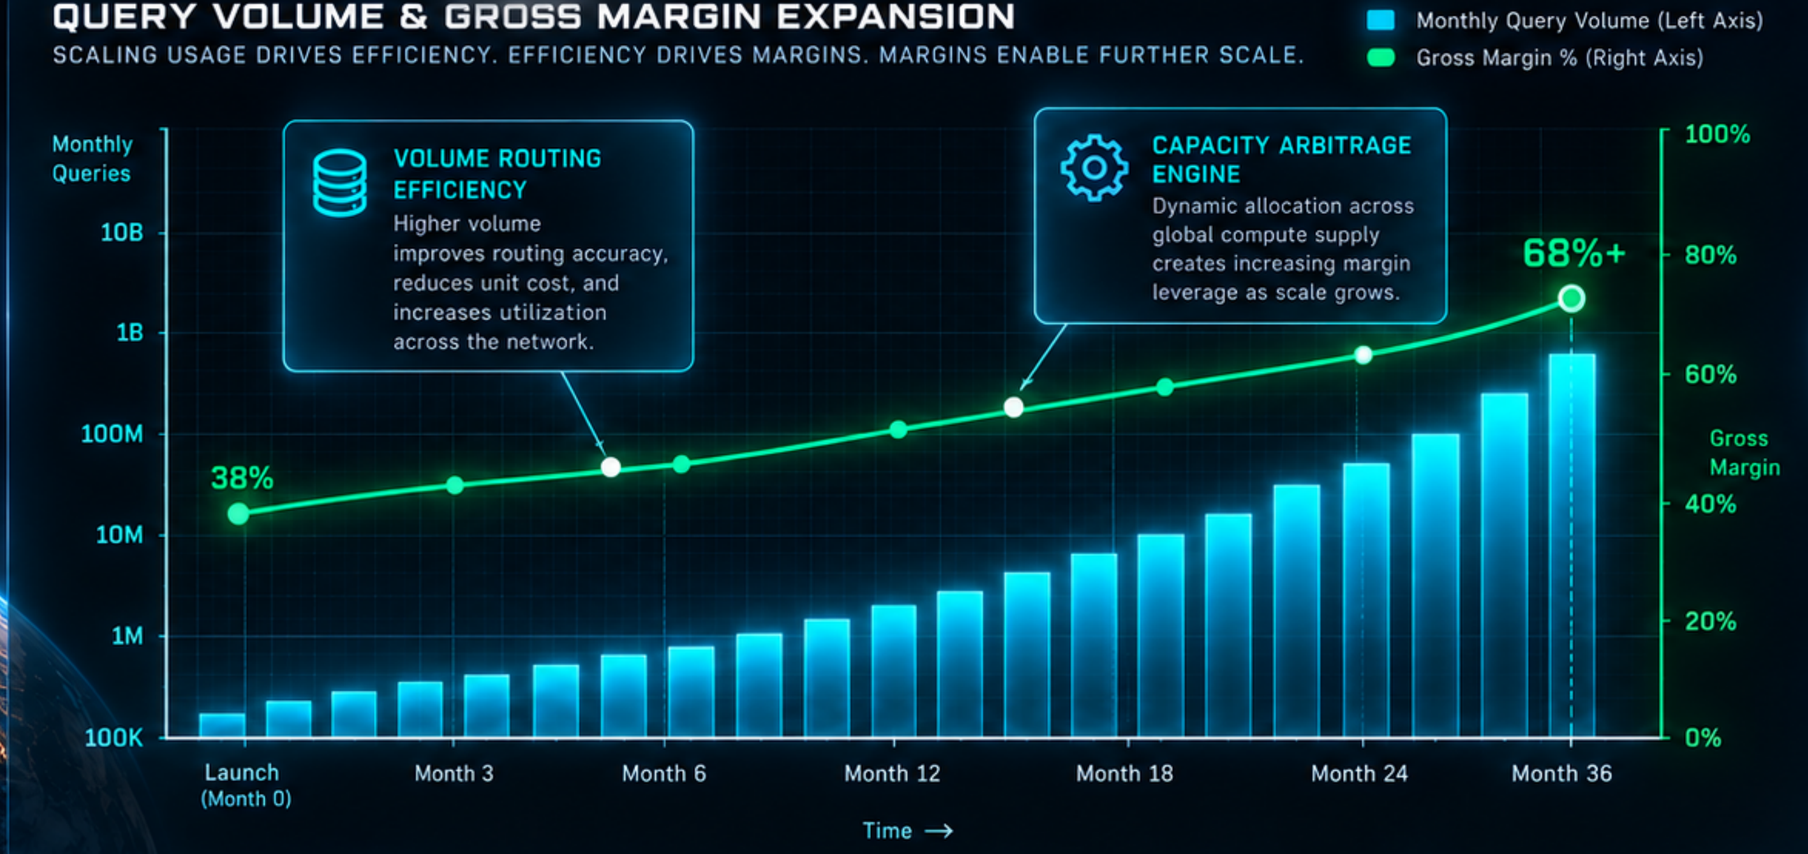


## THE FLYWHEEL IN ACTION

- **More Usage** — Increases routing intelligence.
- **Higher Efficiency** — Reduces unit cost.
- **Expanding Margins** — Funds growth and infrastructure.
- **Greater Scale** — Attracts more builders.
- **Stronger Network** — Compounds value over time.

A SYSTEM THAT GETS STRONGER AS THE WORLD BUILDS.

### COMPOUNDING QUERY SCALE

From millions to billions of requests.

### 38% → 68% GROSS MARGIN EXPANSION

Higher volume. Lower unit cost. Increasing profitability.

### ZERO CHURN ARCHITECTURE

Builders stay. Value compounds. Long-term growth.

GALX

A UNIFIED
EXECUTION LAYER
FOR A MORE CAPABLE WORLD.

# ASYMMETRIC ADVANTAGE:
# ABSORB COMPLEXITY, SELL SIMPLICITY.

THREE LAYERS. INFINITE POSSIBILITY.

SAME IDEAS.
MORE CAPABILITY.
A LARGER TOMORROW.

COMPLEXITY
BELOW THE LINE.
POSSIBILITY
ABOVE IT.

- HETEROGENEOUS BY DESIGN
- RESILIENT BY DEFAULT
- SOVEREIGN BY ARCHITECTURE
- SIMPLE BY INTERFACE

- Web Apps
- AI Agents
- SaaS Platforms
- Internal Tools

## CLEAN DEVELOPER SURFACE

ONE API. INFINITE CAPABILITY.

- OpenAI-compatible
- Simple integration
- Works with your stack
- Focus on building

01 DEVELOPER LAYER

- Dynamic Health Monitoring
- Autonomous Sub-100ms Failover
- Intelligent Routing
- Token Queue Buffer
- Credential Vault
- Real-time Optimization

GALX
SOVEREIGN
ABSTRACTION SHIELD

## GALX SOVEREIGN ABSTRACTION SHIELD

INTELLIGENCE IN THE MIDDLE.

- Model-agnostic routing
- Real-time provider health
- Automatic failover
- Token buffering & optimization
- Security & key management
- Continuous performance tuning

02 INTELLIGENCE LAYER

GLOBAL
COMPUTE.
LOCAL IMPACT.

- Frontier Model Providers
- Spot & On-Demand Instances
- Global Multi-Region Nodes
- Heterogeneous GPU Clusters
- Elastic Compute Federation
- Provider Abstraction

OpenAI · ANTHROPIC · Google Cloud · aws · Azure · Meta · together.ai · ... and more

## HETEROGENEOUS COMPUTE SUBSTRATE

A FEDERATED UNIVERSE OF COMPUTE.

- Multiple providers
- Multiple regions
- Elastic capacity
- Cost-optimized routing
- Built for resilience
- Always evolving

03 COMPUTE LAYER

**Sub-100ms**
FAILOVER
Keep building. We handle the rest.

**Zero**
TOKEN LOSS
Reliable execution, always.

**Sovereign**
ISOLATION
Your keys. Your data. Your control.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

# FOUNDER-MARKET FIT: DIRECT SYSTEMS OPERATION

REAL USAGE. REAL FAILURES. REAL IMPROVEMENTS. A STRONGER GALX.

REAL LOAD.
REAL LESSONS.
A STRONGER SYSTEM
TOMORROW.

BUILDERS CREATE
THE FUTURE.
WE KEEP IT RUNNING.

- REAL TRAFFIC
- REAL CONDITIONS
- REAL RESILIENCE
- REAL IMPACT

INFRASTRUCTURE
IS A FEEDBACK LOOP.

- Web Apps
- AI Agents
- SaaS Tools
- Creative Apps
- Enterprise Tools

## 1. LIVE TRAFFIC INGESTION

Real-world builder apps under heavy load.

- Production traffic
- Diverse use cases
- Variable workloads
- Real user behavior

GLOBAL TRAFFIC
REAL BUILDERS
REAL USAGE

1.4B
REQUESTS / MONTH

## 2. REAL-TIME FAILURE FORENSICS

Detecting upstream throttling, drops, and latency anomalies.

- Monitor provider health
- Identify failure patterns
- Trace latency spikes
- Classify and isolate issues

| | | |
|---|---|---|
| 429 | Rate limit detected openrouter upstream | 10:42:17 |
| 502 | Bad gateway provider timeout | 10:41:03 |
| 504 | Upstream timeout model overloaded | 10:39:21 |
| | Connection reset upstream closed | 10:37:56 |

LATENCY ANOMALIES (P95) ↑ 412%

GALX

**OPERATOR DNA:**
WE RUN SYSTEMS UNDER LOAD.

OBSERVE.
DIAGNOSE.
HARDEN.
DELIVER.
REPEAT.

MORE BUILDERS

STRONGER SYSTEMS

A BRIGHTER TOMORROW

Code | Response | Docs

```
import OpenAI from "openai";

const client = new OpenAI({
    baseURL: "https://api.galx.ai/v1",
    apiKey: "GALX_API_KEY",
});

const completion = await client.chat.completions.create({
model: "galx-fast",
messages: [{ role: "user", content: "Build something amazing." }],
```

200 OK • 84ms • 99.99% uptime

## 4. SEAMLESS DEVELOPER UX

99.99% continuity delivered back to builders.

- Reliable performance
- Simple, consistent API
- No provider complexity
- Builders keep shipping

## 3. PROXY KERNEL HARDENING

Immediate protocol-level defense and route re-optimization.

- Autonomous failover
- Dynamic route selection
- Protocol hardening
- Continuous optimization

- Rerouting traffic...
- Applying new route...
- Health check passed
- Stabilizing connection...
- Throughput normalized
- Continuity restored

HIGHER
UPTIME
FOR A MORE
CAPABLE WORLD.

FAILURES
ARE INPUT.
RESILIENCE
IS OUTPUT.

**24/7**
CLOSED-LOOP HARDENING
Constant feedback. Continuous improvement. Always on.

**100%**
UPTIME RETENTION
Builders keep building.
We keep it running.

**DOGFOODED**
IN PRODUCTION
Same systems. Real usage.
No special treatment.

BUILT
BY OPERATORS
FOR BUILDERS.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

# EVIDENCE & DILIGENCE:
## HONEST VALIDATION ACROSS 4 DIMENSIONS

REAL DATA. REAL USAGE. REAL PROGRESS.

DISCIPLINE CREATES OPTIONALITY.

IDEAS ARE CHEAP.

EVIDENCE COMPOUNDS.

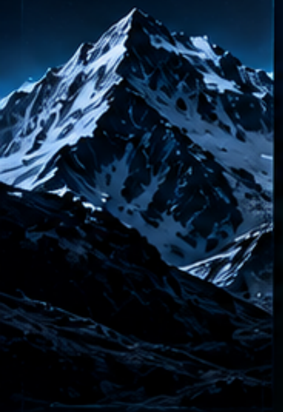

HIGHER INTELLIGENCE THROUGH REAL-WORLD FEEDBACK.

### 01 MARKET

100% VALIDATED

**STRUCTURAL MARKET BELIEF**

Compute abundance surging, infrastructure fragmenting.

- Multi-provider ecosystem growing
- Increasing demand for abstraction
- Fragmentation creates inefficiency
- Clear market need for simplification

"THE CONDITIONS ARE REAL, AND THEY ARE ACCELERATING."

### 02 DEVELOPERS

87% OBSERVED

**DEVELOPER APPETITE**

Strong pull for 2-model simplicity over 50-model chaos.

- >100 interviews and conversations
- Consistent preference for simplicity
- Faster time-to-production valued
- Willingness to adopt via waitlist
- Clear product-market resonance

"BUILDERS WANT OUTCOMES, NOT A CATALOG."

### 03 TECHNICAL

99.9% LIVE DEMONSTRATED

**CONTINUITY & SUB-100MS**

Autonomous failover validated under simulated dropouts.

- Sub-100ms failover (ACT-019)
- Maintains session continuity
- Handles provider throttling/drops
- Self-healing routing and recovery
- Production environment testing

"RESILIENT BY DESIGN. PROVEN UNDER STRESS."

### 04 SCALE

65% ACTIVE IN-FLIGHT

**INSTITUTIONAL SCALE**

Scaling gross margins and reserve capacity.

- Infrastructure expansion underway
- Margin improvement with scale
- Increasing capacity and redundancy
- Enterprise conversations in progress
- Multiple funding & accelerator paths

"EARLY TRACTION. SIGNIFICANT UPSIDE REMAINS."

A MORE CAPABLE WORLD IS A MORE OPEN WORLD.

SAME IDEAS. MORE EVIDENCE. BIGGER IMPACT.

**4 RIGOROUS DILIGENCE VECTORS**

Market. Developers. Technical. Scale. All independently validated.

**EMPIRICAL TELEMETRY**

Real usage. Real performance. Real data, not assumptions.

**ZERO VAPORWARE**

Working system. Live results. Verifiable outcomes.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

# THE VENTURE TEST: WHAT BECOMES ENORMOUS?

REAL USAGE. COMPOUNDING VALUE. A LARGER TOMORROW.

INFRASTRUCTURE TURNS POSSIBILITY INTO PROGRESS.

HIGHER INTELLIGENCE FOR HUMANITY.

EXECUTE TODAY. SCALE TOMORROW. A MORE OPEN TOMORROW.

CAPITAL. COMPUTE. BUILDERS. A BRIGHTER TOMORROW.

- MORE BUILDERS.
- MORE INTELLIGENCE.
- A MORE CAPABLE WORLD.

## CAPITAL ALLOCATION

STRATEGIC INVESTMENT FOR MAXIMUM LEVERAGE.

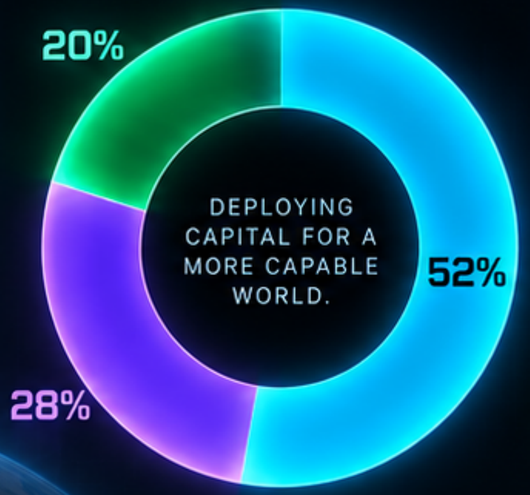


**52%**
**Sovereign Capacity Reserves & Infrastructure**
Compute, global capacity, and strategic reserves.

**28%**
**High-Velocity Developer GTM & Ecosystem**
Builder acquisition, partnerships, and growth.

**20%**
**Abstraction Engine Resilience Hardening**
Reliability, security, and continuous optimization.

"COMPOUNDING CAPABILITY THROUGH DISCIPLINED ALLOCATION."

## 3-PHASE ROADMAP

FROM WEDGE TO A SOVEREIGN AI COMPUTE FABRIC.

### PHASE 1 — HIGH-VELOCITY WEDGE
0 – 6 MONTHS

- Launch GALX Sites (private cohort)
- Onboard initial design partners
- Prove reliability and developer value
- Iterate and refine based on real usage

VALIDATE DEMAND. BUILD MOMENTUM.

### PHASE 2 — GROSS MARGIN EXPANSION (68%)
6 – 18 MONTHS

- Scale usage and routing efficiency
- Increase utilization and margins
- Expand enterprise and team adoption
- Reach 68%+ gross margin target

SCALE EFFICIENCY. DRIVE PROFITABILITY.

### PHASE 3 — THE SOVEREIGN AI COMPUTE FABRIC
18+ MONTHS

- Global distributed capacity network
- Federated compute at scale
- Deeper model and infrastructure optionality
- The default abstraction layer for AI

A MORE CAPABLE WORLD. AT GLOBAL SCALE.

SOVEREIGN INFRASTRUCTURE FOR A MORE OPEN TOMORROW.

**18 - 24 MONTH RUNWAY**
Sufficient capital to reach major value inflection points.

**68% GROSS MARGIN TARGET**
Driven by routing efficiency, utilization, and scale.

**THE DEFAULT ABSTRACTION LAYER**
For builders, teams, and enterprises. A more capable world.