GALX

Sustainable, Abundant AI Access
for Anyone Who Wants to Build

Two Models

<100ms TTFT

99.99% Continuity

Zero Ops

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD

THE STRUCTURAL INEFFICIENCY: THE STATUS QUO TAX

ONE INTERFACE. MANY WORLDS. REAL PROGRESS.

UNSHIELDED STATUS QUO

14 APIs. COUNTLESS PROBLEMS.

INTEGRATION
MAINTENANCE
TROUBLESHOOTING
CONSTANT CHANGES
LOST PRODUCTIVITY

FROM CHAOS TO CLARITY

THE GALX ABSTRACTION SHIELD

ONE INTERFACE. A MORE CAPABLE WORLD.

68%
OPS OVERHEAD
Average engineering time spent on integration, maintenance, and firefighting.

14
FRAGMENTED SDKs
Different APIs. Different rules. A moving target.

<100ms
AUTONOMOUS ROUTING
Intelligent model selection. Optimal performance, automatically.

GALX | UNIFIED INTELLIGENCE REAL PROGRESS.

THE SUFFICIENCY PRINCIPLE:
SUFFICIENT INTELLIGENCE ECONOMICALLY DELIVERED

SAME IDEAS.
MORE EXECUTION.
A MORE CAPABLE WORLD.

TASK COMPLETION vs. MODEL COMPUTE SCALE

Diminishing returns are real. Sufficient intelligence wins.

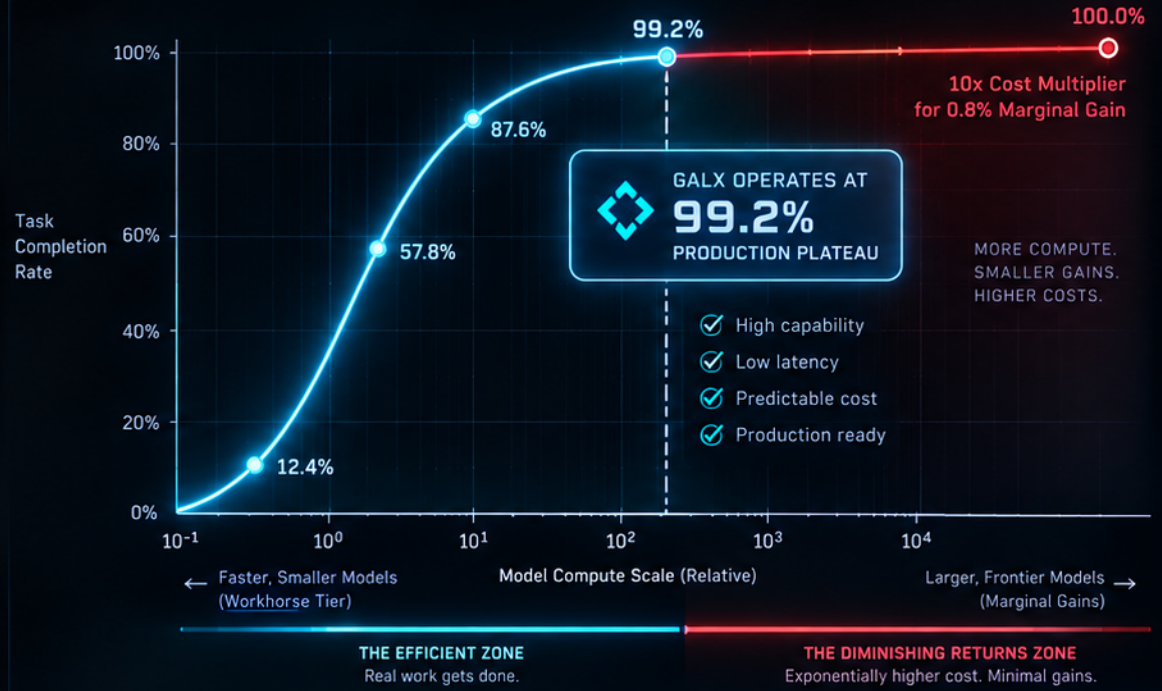

THE VENTURE LANDSCAPE

From Fragile Sprawl to Sovereign Sufficiency.

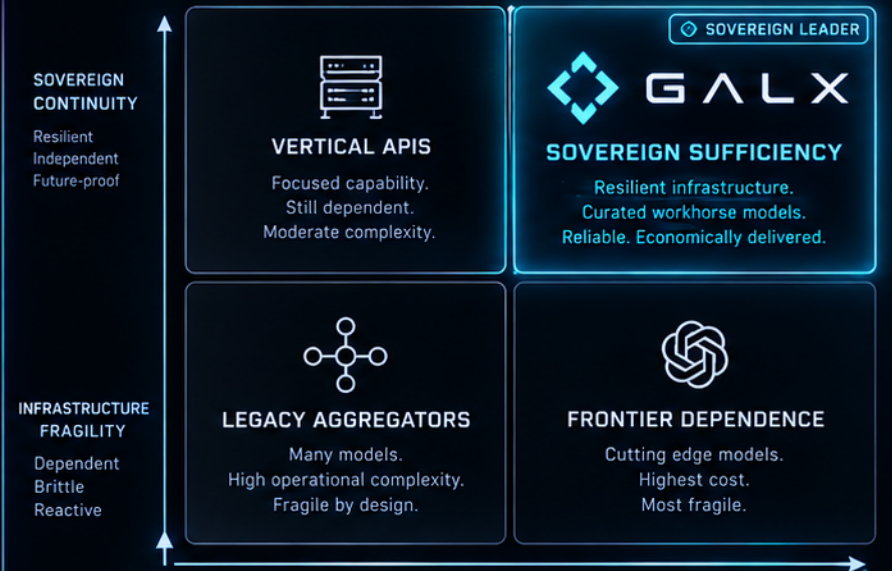

99.2%
TASK PLATEAU
Real-world tasks completed with workhorse models.

10x
COST MULTIPLIER AVOIDED
Frontier-scale costs for only 0.8% marginal gain.

2
CURATED MODELS
Sufficient. Capable. Production ready.
No catalog sprawl.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

RADICAL SIMPLICITY:
TWO MANAGED MODELS. STANDARD API.

LESS COMPLEXITY.
MORE POSSIBILITY.
SAME INTERFACE.

THE RIGHT MODEL FOR THE JOB.
PRODUCTION-READY. ECONOMICALLY DELIVERED.

GALX FAST
SPEED FOR REAL WORK

High-throughput, low-latency inference for everyday tasks. Get to answers, faster.

- **Sub-100ms latency** — Blazing fast responses
- **High throughput** — Built for production scale
- **Micro-cents pricing** — Maximum efficiency
- **Reliable & deterministic** — Production ready

GALX FAST

FAST. EFFICIENT. RELIABLE.

GALX HEAVY
DEPTH FOR COMPLEXITY

Advanced reasoning, long-context understanding, and multi-step agentic workflows.

- **Deep reasoning** — Tackle complex problems
- **High-context synthesis** — Larger context, richer understanding
- **Complex agents** — Multi-step task execution
- **Production-grade reliability** — Same simple interface

GALX HEAVY

DEEPER. SMARTER. FURTHER.

GALX | Code | Response | Logs | Live | cURL

```
curl https://api.galx.ai/v1/chat/completions \
  -H "Content-Type: application/json" \
  -H "Authorization: Bearer $GALX_API_KEY" \
  -d '{
    "model": "galx-fast",
    "messages": [{"role": "user", "content": "Build something amazing."}]
}'
```

Streaming response... 342 tokens/s

```
{
  "id": "galx-9f3c2e1b",
  "object": "chat.completion.chunk",
  "model": "galx-fast",
  "choices": [{
    "delta": {
      "content": "Building",
    }
  }]
}
{  "delta": { "content": " something" } }
{  "delta": { "content": " amazing" } }
{  "delta": { "content": "." } }
```

Status **200 OK**

TTFT **84ms**

Continuity **100%**

Single Unified Endpoint

SAME INTERFACE.
TWO MODELS.
INFINITE POTENTIAL.

1
STANDARD ENDPOINT
One API. Two managed models. No complexity.

<100ms
TTFT
Get to answers, faster. Built for real work.

ZERO
MIGRATION CONFIG
No SDK sprawl. No rewrites. Start building today.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

THE BUILDER WEDGE: HIGH-VELOCITY CREATORS

LESS SETUP. MORE BUILDING. A MORE CAPABLE WORLD.

IDEAS SHIP FASTER HERE.

SAME AMBITION. A SHORTER PATH.

- IDEA TO PROTOTYPE
- PROTOTYPE TO PRODUCT
- BUILD TO IMPACT

INFRASTRUCTURE SHOULD ACCELERATE HUMAN CREATIVITY.

Vercel | supabase | Cloudflare | Next.js

WORKS WITH YOUR STACK.

TIME-TO-PRODUCTION

SAME IDEA. TWO VERY DIFFERENT PATHS.

TRADITIONAL CLOUD / MULTI-VENDOR

Multiple providers.
Multiple roadblocks.
Weeks of friction.

4 to 6 Weeks

SLOWER BUILDS.
HIGHER COSTS.
MORE COMPLEXITY.

GALX PLATFORM

One interface.
Instant access.
Built for builders.

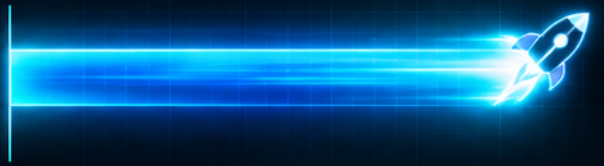

3 Minutes
to First Production Token

SAME IDEAS.
REAL MOMENTUM.

"INFRASTRUCTURE DISAPPEARS. CREATIVITY ACCELERATES."

3-MINUTE ONBOARDING
Go from sign up to production token in 3 minutes.

4–6 WEEKS TIME SAVED
Skip procurement, integration, and infrastructure complexity.

100% FOCUS ON PRODUCT
Spend your time building, not juggling infrastructure.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

WHY THIS CAN BECOME ENORMOUS: THE COMPOUNDING EXECUTION FLYWHEEL

MORE USAGE. HIGHER EFFICIENCY. STRONGER MARGINS. A LARGER GALX.

SAME INTELLIGENCE. A BIGGER TOMORROW.

DEMAND SCALES INTELLIGENCE. EFFICIENCY COMPOUNDS VALUE.

FROM IDEAS TO IMPACT AT GLOBAL SCALE.

A MORE CAPABLE WORLD IS A COMPOUNDING ONE.

QUERY VOLUME & GROSS MARGIN EXPANSION

SCALING USAGE DRIVES EFFICIENCY. EFFICIENCY DRIVES MARGINS. MARGINS ENABLE FURTHER SCALE.

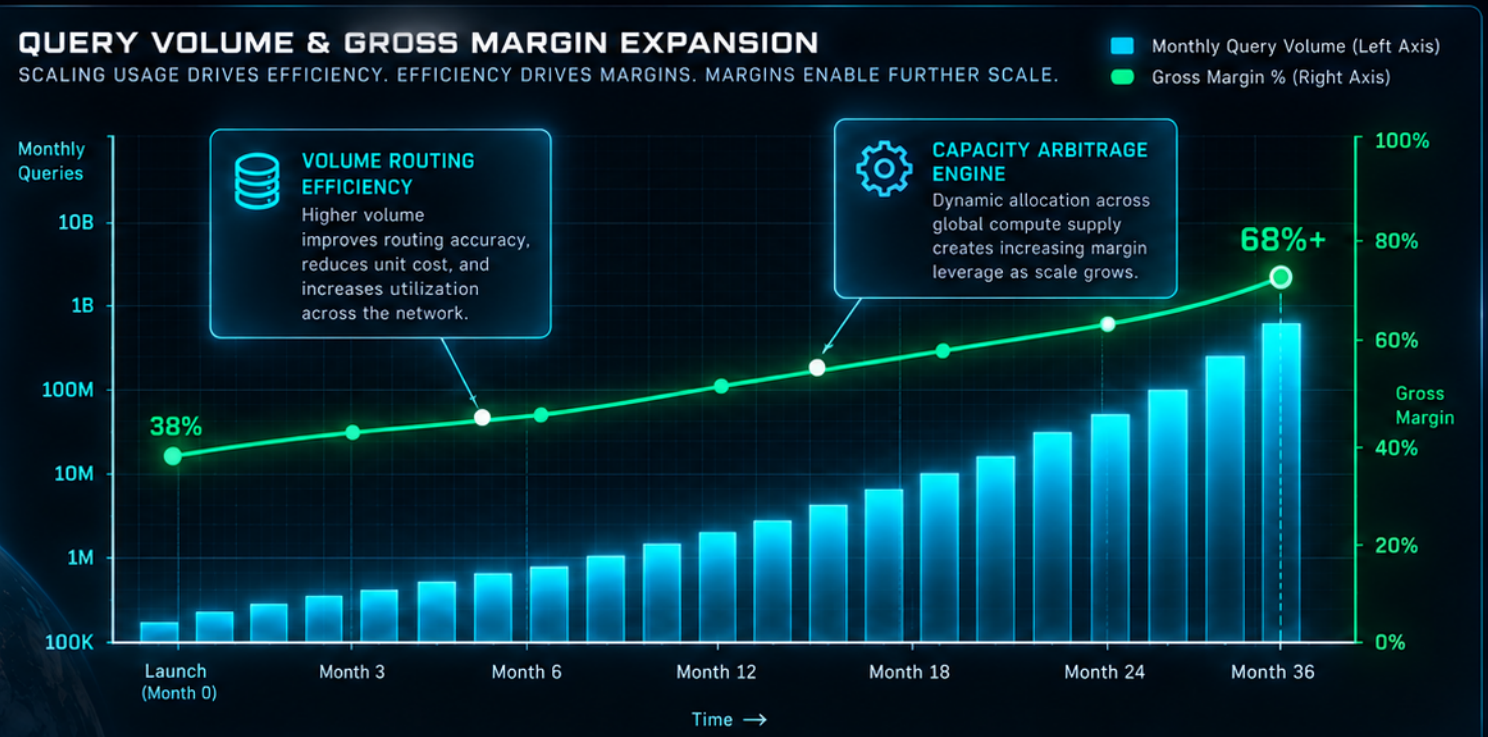

THE FLYWHEEL IN ACTION

- **More Usage** — Increases routing intelligence.
- **Higher Efficiency** — Reduces unit cost.
- **Expanding Margins** — Funds growth and infrastructure.
- **Greater Scale** — Attracts more builders.
- **Stronger Network** — Compounds value over time.

A SYSTEM THAT GETS STRONGER AS THE WORLD BUILDS.

COMPOUNDING QUERY SCALE

From millions to billions of requests.

38% → 68% GROSS MARGIN EXPANSION

Higher volume. Lower unit cost. Increasing profitability.

ZERO CHURN ARCHITECTURE

Builders stay. Value compounds. Long-term growth.

GALX

A UNIFIED
EXECUTION LAYER
FOR A MORE CAPABLE WORLD.

ASYMMETRIC ADVANTAGE:
ABSORB COMPLEXITY, SELL SIMPLICITY.

THREE LAYERS. INFINITE POSSIBILITY.

SAME IDEAS.
MORE CAPABILITY.
A LARGER TOMORROW.

COMPLEXITY
BELOW THE LINE.
POSSIBILITY
ABOVE IT.

- HETEROGENEOUS BY DESIGN
- RESILIENT BY DEFAULT
- SOVEREIGN BY ARCHITECTURE
- SIMPLE BY INTERFACE

- Web Apps
- AI Agents
- SaaS Platforms
- Internal Tools

01 DEVELOPER LAYER

CLEAN DEVELOPER SURFACE

ONE API. INFINITE CAPABILITY.

- OpenAI-compatible
- Simple integration
- Works with your stack
- Focus on building

- Dynamic Health Monitoring
- Autonomous Sub-100ms Failover
- Intelligent Routing
- Token Queue Buffer
- Credential Vault
- Real-time Optimization

GALX
SOVEREIGN
ABSTRACTION SHIELD

02 INTELLIGENCE LAYER

GALX SOVEREIGN ABSTRACTION SHIELD

INTELLIGENCE IN THE MIDDLE.

- Model-agnostic routing
- Real-time provider health
- Automatic failover
- Token buffering & optimization
- Security & key management
- Continuous performance tuning

GLOBAL
COMPUTE.
LOCAL IMPACT.

- Frontier Model Providers
- Spot & On-Demand Instances
- Global Multi-Region Nodes
- Heterogeneous GPU Clusters
- Elastic Compute Federation
- Provider Abstraction

OpenAI · ANTHROPIC · Google Cloud · aws · Azure · Meta · together.ai · ... and more

03 COMPUTE LAYER

HETEROGENEOUS COMPUTE SUBSTRATE

A FEDERATED UNIVERSE OF COMPUTE.

- Multiple providers
- Multiple regions
- Elastic capacity
- Cost-optimized routing
- Built for resilience
- Always evolving

Sub-100ms
FAILOVER
Keep building. We handle the rest.

Zero
TOKEN LOSS
Reliable execution, always.

Sovereign
ISOLATION
Your keys. Your data. Your control.

GALX

A UNIFIED
EXECUTION LAYER
FOR A MORE CAPABLE WORLD.

FOUNDER-MARKET FIT: DIRECT SYSTEMS OPERATION

REAL USAGE. REAL FAILURES. REAL IMPROVEMENTS. A STRONGER GALX.

REAL LOAD.
REAL LESSONS.
A STRONGER SYSTEM
TOMORROW.

BUILDERS CREATE
THE FUTURE.
WE KEEP IT RUNNING.

- REAL TRAFFIC
- REAL CONDITIONS
- REAL RESILIENCE
- REAL IMPACT

INFRASTRUCTURE
IS A FEEDBACK LOOP.

- Web Apps
- AI Agents
- SaaS Tools
- Creative Apps
- Enterprise Tools

GLOBAL TRAFFIC
REAL BUILDERS
REAL USAGE

1.4B
REQUESTS / MONTH

1. LIVE TRAFFIC INGESTION

Real-world builder apps under heavy load.

- Production traffic
- Diverse use cases
- Variable workloads
- Real user behavior

2. REAL-TIME FAILURE FORENSICS

Detecting upstream throttling, drops, and latency anomalies.

- Monitor provider health
- Identify failure patterns
- Trace latency spikes
- Classify and isolate issues

429	Rate limit detected openrouter upstream	10:42:17
502	Bad gateway provider timeout	10:41:03
504	Upstream timeout model overloaded	10:39:21
	Connection reset upstream closed	10:37:56

LATENCY ANOMALIES (P95) ↑ 412%

GALX

OPERATOR DNA:
WE RUN SYSTEMS
UNDER LOAD.

OBSERVE.
DIAGNOSE.
HARDEN.
DELIVER.
REPEAT.

MORE BUILDERS

STRONGER SYSTEMS

A BRIGHTER TOMORROW

4. SEAMLESS DEVELOPER UX

99.99% continuity delivered back to builders.

- Reliable performance
- Simple, consistent API
- No provider complexity
- Builders keep shipping

Code | Response | Docs

```
import OpenAI from "openai";

const client = new OpenAI({
    baseURL: "https://api.galx.ai/v1",
    apiKey: "GALX_API_KEY",
});

const completion = await client.chat.completions.create({
model: "galx-fast",
messages: [{ role: "user", content: "Build something amazing." }],
```

200 OK • 84ms • 99.99% uptime

3. PROXY KERNEL HARDENING

Immediate protocol-level defense and route re-optimization.

- Autonomous failover
- Dynamic route selection
- Protocol hardening
- Continuous optimization

- Rerouting traffic...
- Applying new route...
- Health check passed
- Stabilizing connection...
- Throughput normalized
- Continuity restored

HIGHER
UPTIME
FOR A MORE
CAPABLE WORLD.

FAILURES
ARE INPUT.
RESILIENCE
IS OUTPUT.

24/7
CLOSED-LOOP HARDENING
Constant feedback. Continuous improvement.
Always on.

100%
UPTIME RETENTION
Builders keep building.
We keep it running.

DOGFOODED
IN PRODUCTION
Same systems. Real usage.
No special treatment.

BUILT
BY OPERATORS
FOR BUILDERS.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

EVIDENCE & DILIGENCE:
HONEST VALIDATION ACROSS 4 DIMENSIONS

REAL DATA. REAL USAGE. REAL PROGRESS.

DISCIPLINE CREATES OPTIONALITY.

IDEAS ARE CHEAP.

EVIDENCE COMPOUNDS.

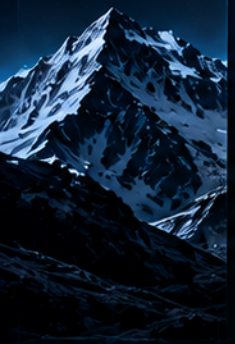

HIGHER INTELLIGENCE THROUGH REAL-WORLD FEEDBACK.

01 MARKET

100% VALIDATED

STRUCTURAL MARKET BELIEF

Compute abundance surging, infrastructure fragmenting.

- Multi-provider ecosystem growing
- Increasing demand for abstraction
- Fragmentation creates inefficiency
- Clear market need for simplification

"THE CONDITIONS ARE REAL, AND THEY ARE ACCELERATING."

02 DEVELOPERS

87% OBSERVED

DEVELOPER APPETITE

Strong pull for 2-model simplicity over 50-model chaos.

- >100 interviews and conversations
- Consistent preference for simplicity
- Faster time-to-production valued
- Willingness to adopt via waitlist
- Clear product-market resonance

"BUILDERS WANT OUTCOMES, NOT A CATALOG."

03 TECHNICAL

99.9% LIVE DEMONSTRATED

CONTINUITY & SUB-100MS

Autonomous failover validated under simulated dropouts.

- Sub-100ms failover (ACT-019)
- Maintains session continuity
- Handles provider throttling/drops
- Self-healing routing and recovery
- Production environment testing

"RESILIENT BY DESIGN. PROVEN UNDER STRESS."

04 SCALE

65% ACTIVE IN-FLIGHT

INSTITUTIONAL SCALE

Scaling gross margins and reserve capacity.

- Infrastructure expansion underway
- Margin improvement with scale
- Increasing capacity and redundancy
- Enterprise conversations in progress
- Multiple funding & accelerator paths

"EARLY TRACTION. SIGNIFICANT UPSIDE REMAINS."

A MORE CAPABLE WORLD IS A MORE OPEN WORLD.

SAME IDEAS. MORE EVIDENCE. BIGGER IMPACT.

4 RIGOROUS DILIGENCE VECTORS

Market. Developers. Technical. Scale. All independently validated.

EMPIRICAL TELEMETRY

Real usage. Real performance. Real data, not assumptions.

ZERO VAPORWARE

Working system. Live results. Verifiable outcomes.

GALX | A UNIFIED EXECUTION LAYER FOR A MORE CAPABLE WORLD.

THE VENTURE TEST: WHAT BECOMES ENORMOUS?

REAL USAGE. COMPOUNDING VALUE. A LARGER TOMORROW.

INFRASTRUCTURE TURNS POSSIBILITY INTO PROGRESS.

HIGHER INTELLIGENCE FOR HUMANITY.

CAPITAL. COMPUTE. BUILDERS. A BRIGHTER TOMORROW.

- MORE BUILDERS.
- MORE INTELLIGENCE.
- A MORE CAPABLE WORLD.

CAPITAL ALLOCATION

STRATEGIC INVESTMENT FOR MAXIMUM LEVERAGE.

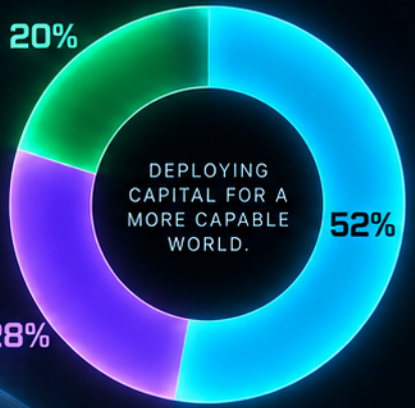

52%
Sovereign Capacity Reserves & Infrastructure
Compute, global capacity, and strategic reserves.

28%
High-Velocity Developer GTM & Ecosystem
Builder acquisition, partnerships, and growth.

20%
Abstraction Engine Resilience Hardening
Reliability, security, and continuous optimization.

"COMPOUNDING CAPABILITY THROUGH DISCIPLINED ALLOCATION."

3-PHASE ROADMAP

FROM WEDGE TO A SOVEREIGN AI COMPUTE FABRIC.

EXECUTE TODAY. SCALE TOMORROW. A MORE OPEN TOMORROW.

PHASE 1
HIGH-VELOCITY WEDGE
0 – 6 MONTHS

- Launch GALX Sites (private cohort)
- Onboard initial design partners
- Prove reliability and developer value
- Iterate and refine based on real usage

VALIDATE DEMAND. BUILD MOMENTUM.

PHASE 2
GROSS MARGIN EXPANSION (68%)
6 – 18 MONTHS

- Scale usage and routing efficiency
- Increase utilization and margins
- Expand enterprise and team adoption
- Reach 68%+ gross margin target

SCALE EFFICIENCY. DRIVE PROFITABILITY.

PHASE 3
THE SOVEREIGN AI COMPUTE FABRIC
18+ MONTHS

- Global distributed capacity network
- Federated compute at scale
- Deeper model and infrastructure optionality
- The default abstraction layer for AI

A MORE CAPABLE WORLD. AT GLOBAL SCALE.

SOVEREIGN INFRASTRUCTURE FOR A MORE OPEN TOMORROW.

18 - 24 MONTH RUNWAY
Sufficient capital to reach major value inflection points.

68% GROSS MARGIN TARGET
Driven by routing efficiency, utilization, and scale.

THE DEFAULT ABSTRACTION LAYER
For builders, teams, and enterprises. A more capable world.